

Input

Persuader turn + context

One member: transformer recipe runs 7 times

Pretrained transformer
RoBERTa,
ELECTRA

Role embedding
who is speaking

Attention pooling
one vector z

Category head
11, main

Strategy head
41, auxiliary

asymmetric loss on both heads

$\times 7$

TF-IDF vectorizer

shared, not a member

Pool of 19 members

2 zero-shot

similarity to label definitions

7 fine-tuned encoders
the only members that train

z

8 few-shot companions

nearest centroid on z or TF-IDF

2 TF-IDF classifiers

LR and XGBoost

19 probabilities for every (turn, category) pair

Decision layer

Ensemble search

50 to 70 combinars,
grouped CV

Calibrated blend
selected, 0.572

XGBoost stacker
tested, 0.551

Decision rule
per label, 1 of 9

Label set
over 11 categories

val macro-F1